

EMAIL SPAM DETECTION WITH MACHINE LEARNING

This notebook was created by DAJAH VINCENT (<https://www.linkedin.com/in/dajahvincent/>)

Importing Data Manipulation Libraries

In [2]: *#Importing data manipulation and visualization libraries*

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

In [3]: *#Reading the CSV file containing the datasets*

```
spam_mail = pd.read_csv('spam.csv', encoding = 'latin')
```

In [4]: *#Viewing the first 5 and last 5 rows from the dataset*

```
spam_mail
```

Out[4]:

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN
...
5567	spam	This is the 2nd time we have tried 2 contact u...	NaN	NaN	NaN
5568	ham	Will Ì_b going to esplanade fr home?	NaN	NaN	NaN
5569	ham	Pity, * was in mood for that. So...any other s...	NaN	NaN	NaN
5570	ham	The guy did some bitching but I acted like i'd...	NaN	NaN	NaN
5571	ham	Rofl. Its true to its name	NaN	NaN	NaN

5572 rows × 5 columns

In [5]:

```
spam_mail.columns
```

Out[5]: Index(['v1', 'v2', 'Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], dtype='object')

In [6]:

```
spam_mail.isna().sum()
```

Out[6]:

v1	0
v2	0
Unnamed: 2	5522
Unnamed: 3	5560
Unnamed: 4	5566
dtype:	int64

In [7]:

```
spam_mail.describe()
```

Out[7]:

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
count	5572	5572	50	12	6
unique	2	5169	43	10	5
top	ham	Sorry, I'll call later	bt not his girlfrnd... G o o d n i g h t . . . @ "	MK17 92H. 450Ppw 16"	GNT:-)"
freq	4825	30	3	2	2

In [8]:

```
spam_mail = spam_mail.drop(columns = ["Unnamed: 2","Unnamed: 3","Unnamed: 4"], axis = 1, errors = 'ignore')
```

In [9]: spam_mail

Out[9]:

	v1	v2
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...
...
5567	spam	This is the 2nd time we have tried 2 contact u...
5568	ham	Will i_b going to esplanade fr home?
5569	ham	Pity, * was in mood for that. So...any other s...
5570	ham	The guy did some bitching but I acted like i'd...
5571	ham	Rofl. Its true to its name

5572 rows × 2 columns

```
In [10]: df = spam_mail.rename(columns = {'v1': 'Class', 'v2': 'Text'})
df
```

Out[10]:

	Class	Text
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...
...
5567	spam	This is the 2nd time we have tried 2 contact u...
5568	ham	Will l_b going to esplanade fr home?
5569	ham	Pity, * was in mood for that. So...any other s...
5570	ham	The guy did some bitching but I acted like i'd...
5571	ham	Rofl. Its true to its name

5572 rows × 2 columns

```
In [11]: df.duplicated().sum()
```

Out[11]: 403

```
In [12]: df = df.drop_duplicates()
df
```

Out[12]:

	Class	Text
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...
...
5567	spam	This is the 2nd time we have tried 2 contact u...
5568	ham	Will l_b going to esplanade fr home?
5569	ham	Pity, * was in mood for that. So...any other s...
5570	ham	The guy did some bitching but I acted like i'd...
5571	ham	Rofl. Its true to its name

5169 rows × 2 columns

```
In [13]: #df['Class'].replace({"ham": 0, "spam": 1}, inplace=True)
#df
```

```
In [14]: df.duplicated().sum()
```

Out[14]: 0

Training & Evaluation

- Importing all Necessary Machine Learning Models & Libraries

In [15]: *# Import necessary libraries*

```
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
```

- Using Support Vector Machine Model

In [16]:

```
# Assigning the feature and target columns from the datasets
X = df['Text'] # Feature
y = df['Class'] # Target

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)

# Convert text data into TF-IDF vectors
vectorizer = TfidfVectorizer(max_features=1000) # You can adjust the max_features
X_train_tfidf = vectorizer.fit_transform(X_train)
X_test_tfidf = vectorizer.transform(X_test)

# Train an SVM model
clf = SVC(kernel='linear', C=1.0, probability=True)
clf.fit(X_train_tfidf, y_train)

# Make predictions on the test data
y_pred = clf.predict(X_test_tfidf)

# Evaluate the model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```


Accuracy: 0.9874274661508704

Classification Report:

	precision	recall	f1-score	support
ham	0.99	1.00	0.99	889
spam	0.98	0.93	0.95	145
accuracy			0.99	1034
macro avg	0.98	0.96	0.97	1034
weighted avg	0.99	0.99	0.99	1034

```
In [17]: # Assuming you already have 'y_test' (true labels) and 'y_pred' (predicted labels)
conf_matrix = confusion_matrix(y_test, y_pred)

# Create a heatmap for the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix - Support Vector Machines')
plt.show()
```



- Using Random Forest Classifier

```
In [18]: # Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(df['Text'], df['Class'], test_size = 0.2, random_state = 42)

# Create a TF-IDF vectorizer
vectorizer = TfidfVectorizer(max_features = 1000) # The max_features can be adjusted as deem fit by users

# Transform text data into numerical features
X_train_tfidf = vectorizer.fit_transform(X_train)
X_test_tfidf = vectorizer.transform(X_test)

# Initialize and train the Random Forest classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
rf_classifier.fit(X_train_tfidf, y_train)

# Predictions on the test set
y_pred = rf_classifier.predict(X_test_tfidf)
```

```
In [19]: # Accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")

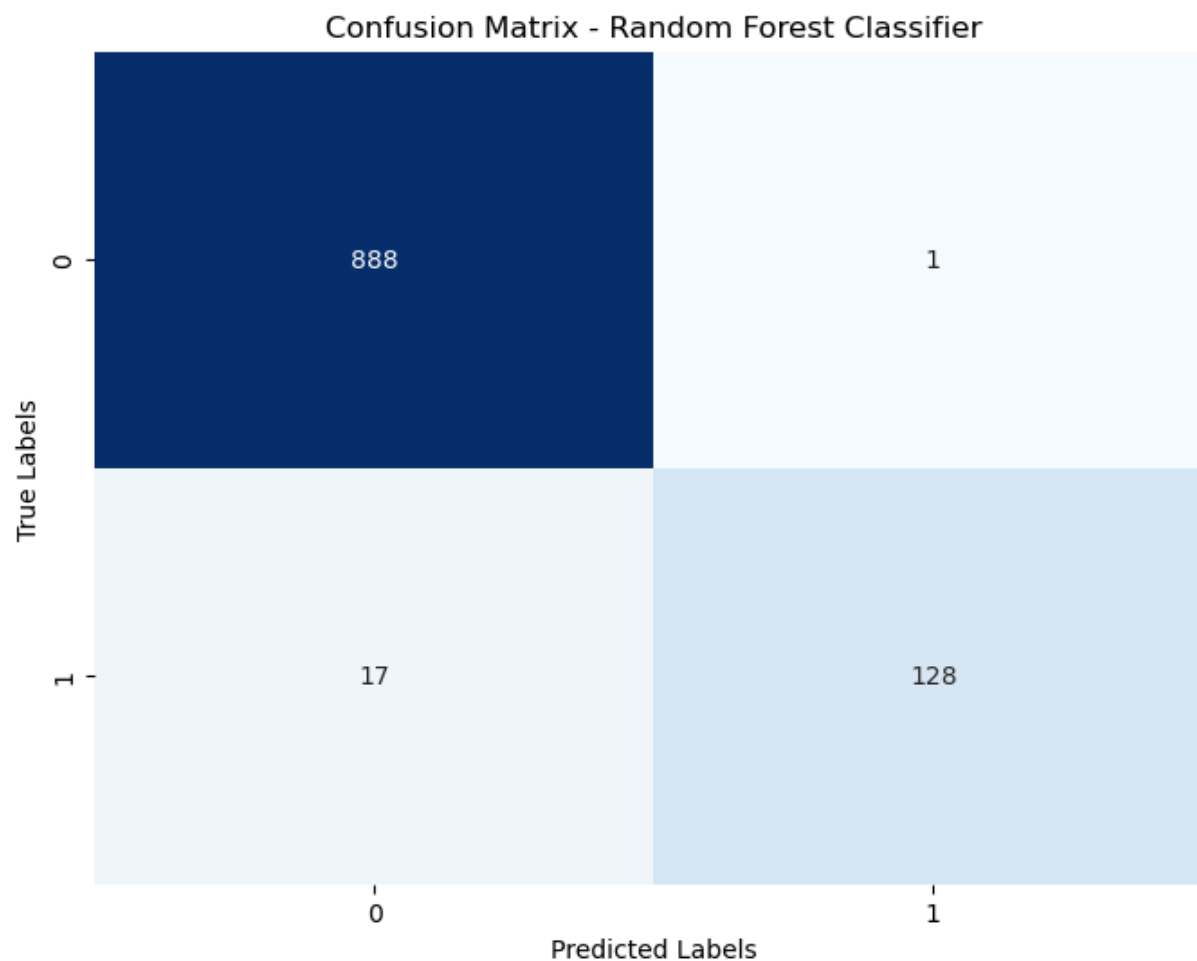
# Classification report (precision, recall, F1-score)
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9826

	precision	recall	f1-score	support
ham	0.98	1.00	0.99	889
spam	0.99	0.88	0.93	145
accuracy			0.98	1034
macro avg	0.99	0.94	0.96	1034
weighted avg	0.98	0.98	0.98	1034

```
In [20]: # Assuming you already have 'y_test' (true labels) and 'y_pred' (predicted labels)
conf_matrix = confusion_matrix(y_test, y_pred)

# Create a heatmap for the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix - Random Forest Classifier')
plt.show()
```



- Using Logistics Regression Model

```
In [21]: # Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(df['Text'], df['Class'], test_size=0.2, random_state=42)

# Create a TF-IDF vectorizer
vectorizer = TfidfVectorizer(max_features = 4000) # Adjust max_features as needed

# Transform text data into numerical features
X_train_tfidf = vectorizer.fit_transform(X_train)
X_test_tfidf = vectorizer.transform(X_test)

# Initialize and train the logistic regression model
logreg_model = LogisticRegression()
logreg_model.fit(X_train_tfidf, y_train)

# Predictions on the test set
y_pred = logreg_model.predict(X_test_tfidf)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")


print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.9710
```

	precision	recall	f1-score	support
ham	0.97	1.00	0.98	889
spam	0.97	0.82	0.89	145
accuracy			0.97	1034
macro avg	0.97	0.91	0.94	1034
weighted avg	0.97	0.97	0.97	1034

```
In [22]: # Assuming you already have 'y_test' (true labels) and 'y_pred' (predicted labels)
conf_matrix = confusion_matrix(y_test, y_pred)

# Create a heatmap for the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix - Logistics Regression')
plt.show()
```



- Using Naive Bayes - MultinomialNB Model

```
In [23]: # Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(df['Text'], df['Class'], test_size = 0.2, random_state = 42)

# Create a TF-IDF vectorizer
vectorizer = TfidfVectorizer(max_features=1000) # Adjust max_features as needed

# Transform text data into numerical features
X_train_tfidf = vectorizer.fit_transform(X_train)
X_test_tfidf = vectorizer.transform(X_test)

# Initialize and train the Naive Bayes classifier
nb_classifier = MultinomialNB()
nb_classifier.fit(X_train_tfidf, y_train)

# Predictions on the test set
y_pred = nb_classifier.predict(X_test_tfidf)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")

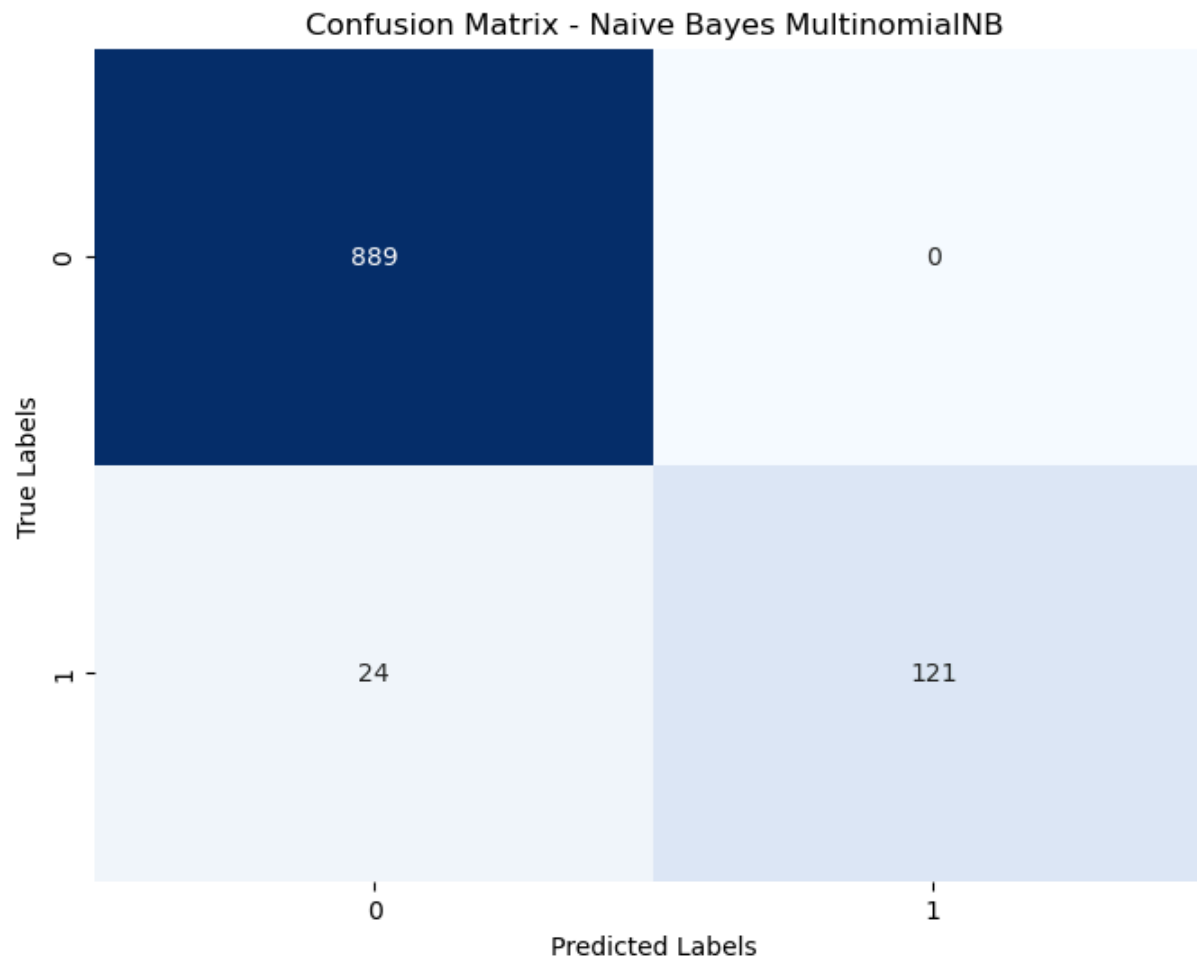
print(classification_report(y_test, y_pred))
```

Accuracy: 0.9768

	precision	recall	f1-score	support
ham	0.97	1.00	0.99	889
spam	1.00	0.83	0.91	145
accuracy			0.98	1034
macro avg	0.99	0.92	0.95	1034
weighted avg	0.98	0.98	0.98	1034

```
In [24]: # Compute the confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix - Naive Bayes MultinomialNB')
plt.show()
```



Testing the Model Using Unseen Text & Text From the Training Datasets

- Using Unseen Text to Test the Prediction of the Model

```
In [25]: #Using unseen text to test the models prediction

# Example text for prediction
new_text = input("Please type a word here: ")

# Transform the new text into numerical features
new_text_tfidf = vectorizer.transform([new_text])

# Make predictions
predicted_class = nb_classifier.predict(new_text_tfidf)

print(f"Predicted class: {predicted_class[0]}")
```

Please type a word here: ghgf hfdhd
Predicted class: ham

The model was tested with a completely different text that is not on the training data set and it with key words like "Won", "20000 dollar" and a hyperlink "www.dajah.com" (<http://www.dajah.com>) and the model correctly predicted the text as Spam

- Randomly Sampling Text from the Dataset to Test the Prediction of the Model

```
In [27]: import random

# Sample 10 random indices from the test set
sample_indices = random.sample(range(len(X_test)), 10)

# Initialize lists to store results
sample_texts = []
true_labels = []
predicted_labels = []

for idx in sample_indices:
    sample_text = X_test.iloc[idx]
    sample_text_tfidf = vectorizer.transform([sample_text])
    predicted_class = nb_classifier.predict(sample_text_tfidf)[0]
    true_label = y_test.iloc[idx]

    sample_texts.append(sample_text)
    true_labels.append(true_label)
    predicted_labels.append(predicted_class)

# Create a DataFrame to display the results
results_df = pd.DataFrame({
    'Sample Text': sample_texts,
    'True Label': true_labels,
    'Predicted Label': predicted_labels
})

# Display the results
print(results_df)
```


	Sample Text	True Label	\
0	Sometimes Heart Remembrs someone Very much... ...	ham	
1	Ok lor but not too early. Me still having proj...	ham	
2	* Was thinking about chuckin ur red green n bl...	ham	
3	Pick you up bout 7.30ish? What time are and t...	ham	
4	Aight will do, thanks again for comin out	ham	
5	Cool, I'll text you when I'm on the way	ham	
6	How's it feel? Mr. Your not my real Valentine ...	ham	
7	I'm not sure if its still available though	ham	
8	Thts wat Wright Brother did to fly..	ham	
9	That seems unnecessarily affectionate	ham	

	Predicted Label
0	ham
1	ham
2	ham
3	ham
4	ham
5	ham
6	ham
7	ham
8	ham
9	ham